

An overview of Recurrent Neural Networks

Jeremy Watt

Features and supervised learning

- Some data types have **independent features**

Features and supervised learning

- Some data types have **independent features**
- **Order in which we feed them in doesn't matter**

Race	Gender	Diabetes	Weight	Readmitted
caucasian	male	No	188	No

Features and supervised learning


- Some data types have **spatially correlated features**

Features and supervised learning

- Some data types have **spatially correlated features**
- **Order in which we feed them in definitely matters**

Features and supervised learning

- Some data types have **spatially correlated features**
- **Order in which we feed them in definitely matters**
- Convolutional modeling can leverage this



Features and supervised learning


- Some data types have **temporarily ordered features**

Features and supervised learning

- Some data types have **temporarily ordered features**
- **Order in which we feed them in definitely matters**


Features and supervised learning


- Some data types have **temporarily ordered features**
- **Order in which we feed them in definitely matters**
- **Recursive modeling can leverage this** (e.g., RNNs)





time







Popular problems with
ordered sequential I/O

Time series prediction

- predict future values of a time series
- **input:** ordered sequence of past series values
- **output:** ordered sequence of future series values



series shown interpolated



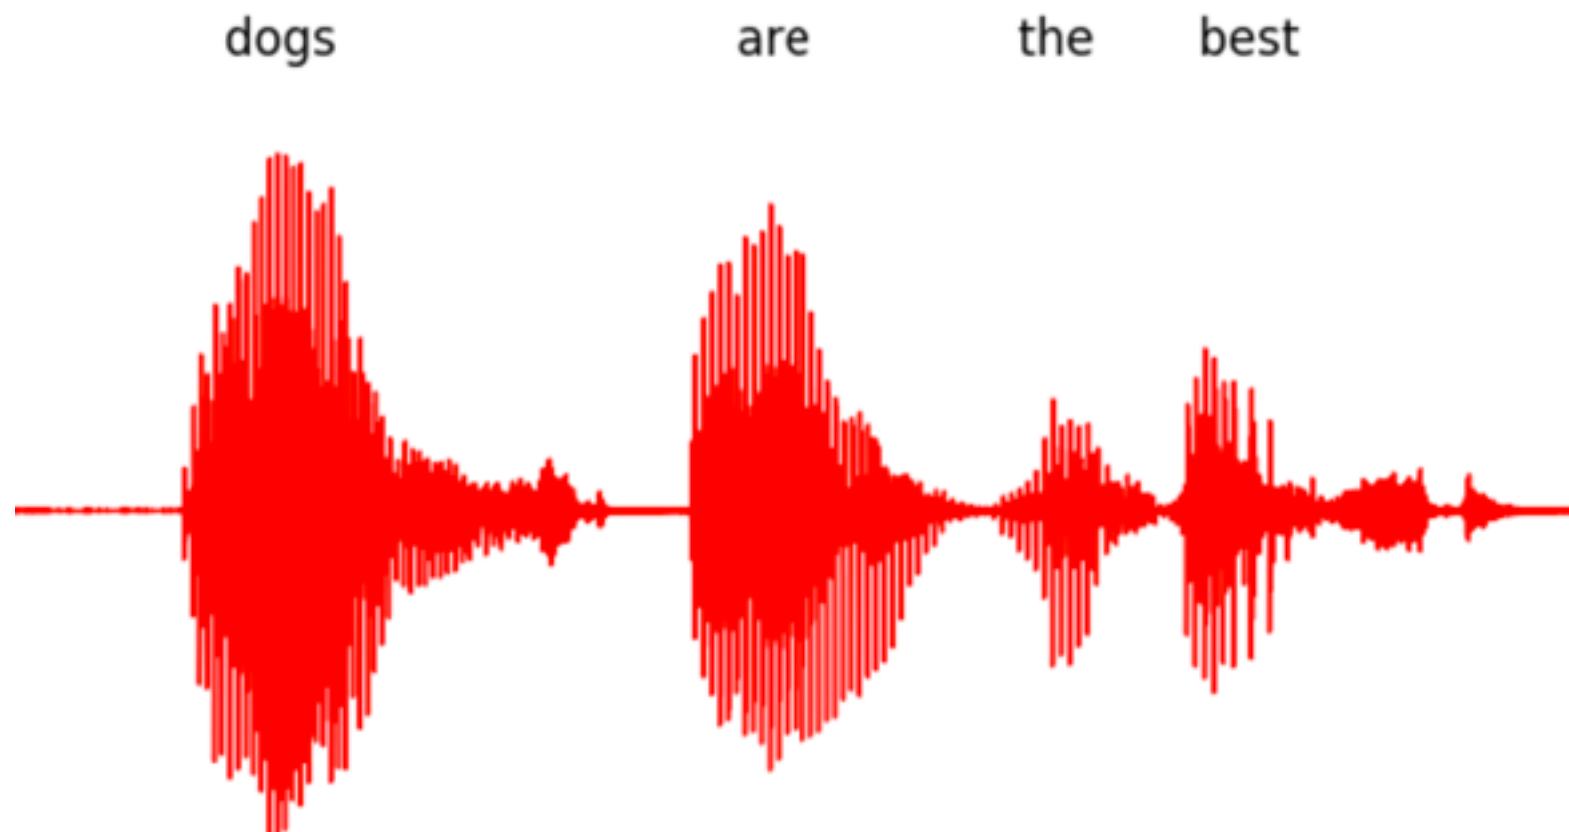
series really looks like this

Machine translation

- Translate language X into language Y
- **input:** ordered sequence of words (in X)
- **output:** ordered sequence of words (in Y)
 - “I do not like cats” —> “Los gatos me caen mal”

Speech recognition

- Speech to text
- **input:** ordered sequence of raw audio
- **output:** ordered sequence of words




Text generation

- generate valid but wacky text automatically
- **input:** ordered sequence of characters (training text corpus)
- **output:** ordered sequence of characters
- Generate wacky sentences with this academic RNN text generator
- Various twitter bots that tweet automatically generated text like this one.
- NanoGenMo annual contest to automatically produce a 50,000+ novel automatically
- Robot Shakespear a text generator that automatically produces Shakespear-esk sentences


Ingesting sequential I/O data
for supervised learning

Time series prediction

- Sequence of P (floating point) numbers: $\langle s_0, s_1, s_2, \dots, s_P \rangle$




series shown interpolated




series really looks like this

Time series prediction

- Sequence of P (floating point) numbers: $\langle s_0, s_1, s_2, \dots, s_P \rangle$
- s_p : the p th value




series shown interpolated




series really looks like this

Time series prediction

- Sequence of P (floating point) numbers: $\langle s_0, s_1, s_2, \dots, s_P \rangle$
- s_p : the p th value
- **regression problem**- I/O are windowed subsequences




series shown interpolated





series really looks like this


Time series prediction


- Sequence of P (floating point) numbers: $\langle s_0, s_1, s_2, \dots, s_P \rangle$
- s_p : the p th value
- - **regression problem**- I/O are windowed subsequences
 - so need training and testing sets
 - lets see how both are formed, start with training












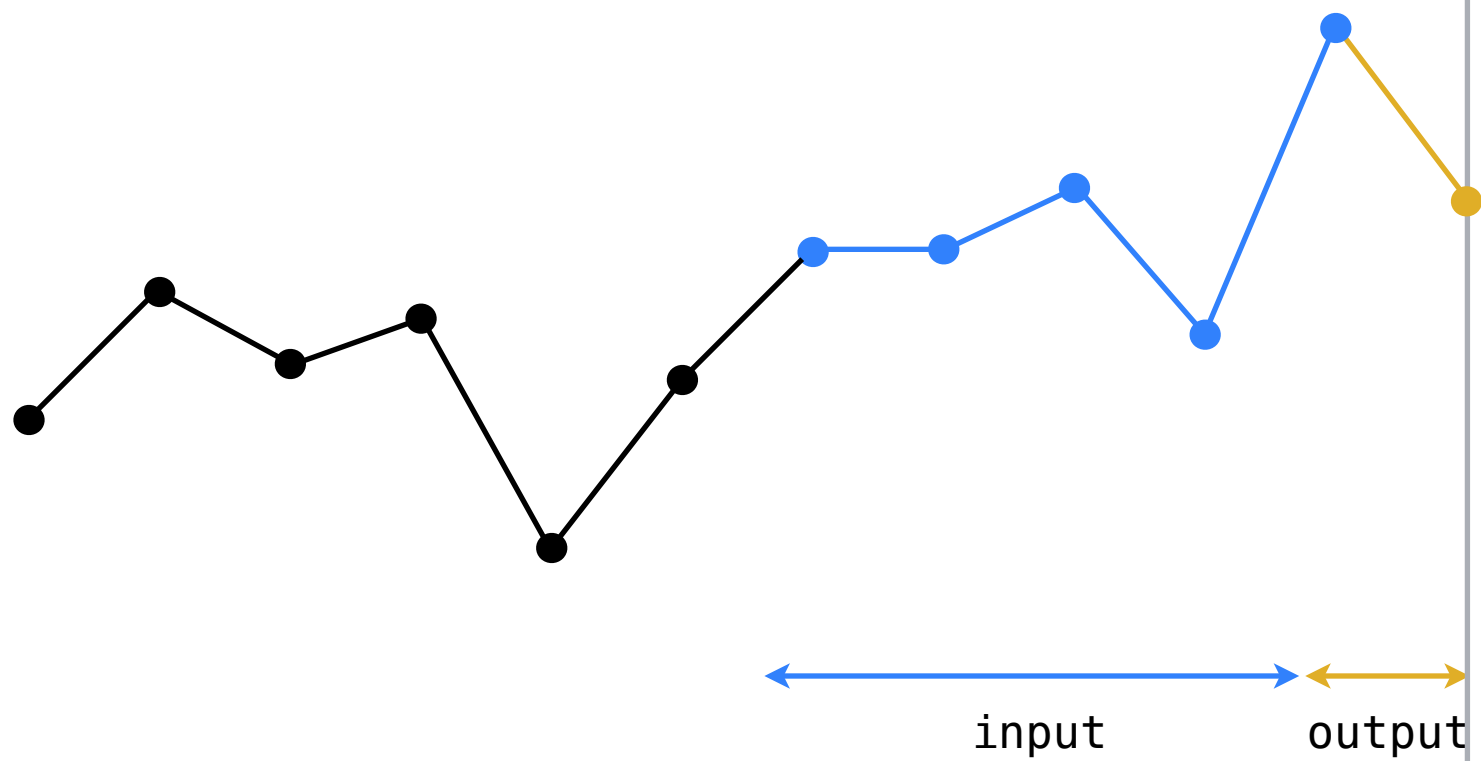


input output












training

testing


One more time - with notation









Input	Output
$\langle s_0, s_1, s_2, s_3 \rangle$	s_4




Input	Output
$\langle s_0, s_1, s_2, s_3 \rangle$	s_4
$\langle s_1, s_2, s_3, s_4 \rangle$	s_5
\vdots	\vdots




Input	Output
$\langle s_0, s_1, s_2, s_3 \rangle$	s_4
$\langle s_1, s_2, s_3, s_4 \rangle$	s_5
\vdots	\vdots




Input	Output
$\langle s_0, s_1, s_2, s_3 \rangle$	s_4
$\langle s_1, s_2, s_3, s_4 \rangle$	s_5
\vdots	\vdots



Input	Output
$\langle s_0, s_1, s_2, s_3 \rangle$	s_4
$\langle s_1, s_2, s_3, s_4 \rangle$	s_5
\vdots	\vdots




Input	Output
$\langle s_0, s_1, s_2, s_3 \rangle$	s_4
$\langle s_1, s_2, s_3, s_4 \rangle$	s_5
\vdots	\vdots
$\langle s_{P-4}, s_{P-3}, s_{P-2}, s_{P-1} \rangle$	s_P



training

testing


Input	Output
$\langle s_0, s_1, s_2, s_3 \rangle$	s_4
$\langle s_1, s_2, s_3, s_4 \rangle$	s_5
\vdots	\vdots
$\langle s_{P-4}, s_{P-3}, s_{P-2}, s_{P-1} \rangle$	s_P



← input →


training

testing



training


testing



training


testing

input



training


testing



training


input

testing



training


testing



input


training

testing



training


testing




input

training

testing






← input →

training

testing


One more time - with notation



training

testing

Input	Output
$\langle s_{P-3}, s_{P-2}, s_{P-1}, s_P \rangle$	\hat{s}_{P+1}




← input →

training

testing


Input	Output
$\langle s_{P-3}, s_{P-2}, s_{P-1}, s_P \rangle$	\hat{s}_{P+1}



training

testing


Input	Output
$\langle s_{P-3}, s_{P-2}, s_{P-1}, s_P \rangle$	\hat{s}_{P+1}
$\langle s_{P-2}, s_{P-1}, s_P, \hat{s}_{P+1} \rangle$	\hat{s}_{P+2}




training

testing

Input	Output
$\langle s_{P-3}, s_{P-2}, s_{P-1}, s_P \rangle$	\hat{s}_{P+1}
$\langle s_{P-2}, s_{P-1}, s_P, \hat{s}_{P+1} \rangle$	\hat{s}_{P+2}




Input	Output
$\langle s_{P-3}, s_{P-2}, s_{P-1}, s_P \rangle$	\hat{s}_{P+1}
$\langle s_{P-2}, s_{P-1}, s_P, \hat{s}_{P+1} \rangle$	\hat{s}_{P+2}
$\langle s_{P-1}, s_P, \hat{s}_{P+1}, \hat{s}_{P+2} \rangle$	\hat{s}_{P+3}




training

testing

Input	Output
$\langle s_{P-3}, s_{P-2}, s_{P-1}, s_P \rangle$	\hat{s}_{P+1}
$\langle s_{P-2}, s_{P-1}, s_P, \hat{s}_{P+1} \rangle$	\hat{s}_{P+2}
$\langle s_{P-1}, s_P, \hat{s}_{P+1}, \hat{s}_{P+2} \rangle$	\hat{s}_{P+3}




Input	Output
$\langle s_{P-3}, s_{P-2}, s_{P-1}, s_P \rangle$	\hat{s}_{P+1}
$\langle s_{P-2}, s_{P-1}, s_P, \hat{s}_{P+1} \rangle$	\hat{s}_{P+2}
$\langle s_{P-1}, s_P, \hat{s}_{P+1}, \hat{s}_{P+2} \rangle$	\hat{s}_{P+3}
$\langle s_P, \hat{s}_{P+1}, \hat{s}_{P+2}, \hat{s}_{P+3} \rangle$	\hat{s}_{P+4}




training

testing

Input	Output
$\langle s_{P-3}, s_{P-2}, s_{P-1}, s_P \rangle$	\hat{s}_{P+1}
$\langle s_{P-2}, s_{P-1}, s_P, \hat{s}_{P+1} \rangle$	\hat{s}_{P+2}
$\langle s_{P-1}, s_P, \hat{s}_{P+1}, \hat{s}_{P+2} \rangle$	\hat{s}_{P+3}
$\langle s_P, \hat{s}_{P+1}, \hat{s}_{P+2}, \hat{s}_{P+3} \rangle$	\hat{s}_{P+4}



Input	Output
$\langle s_{P-3}, s_{P-2}, s_{P-1}, s_P \rangle$	\hat{s}_{P+1}
$\langle s_{P-2}, s_{P-1}, s_P, \hat{s}_{P+1} \rangle$	\hat{s}_{P+2}
$\langle s_{P-1}, s_P, \hat{s}_{P+1}, \hat{s}_{P+2} \rangle$	\hat{s}_{P+3}
$\langle s_P, \hat{s}_{P+1}, \hat{s}_{P+2}, \hat{s}_{P+3} \rangle$	\hat{s}_{P+4}
$\langle \hat{s}_{P+1}, \hat{s}_{P+2}, \hat{s}_{P+3}, \hat{s}_{P+4} \rangle$	\hat{s}_{P+5}



training

testing

here we illustrated with
window size $T = 4$

Input	Output
$\langle s_{P-3}, s_{P-2}, s_{P-1}, s_P \rangle$	\hat{s}_{P+1}
$\langle s_{P-2}, s_{P-1}, s_P, \hat{s}_{P+1} \rangle$	\hat{s}_{P+2}
$\langle s_{P-1}, s_P, \hat{s}_{P+1}, \hat{s}_{P+2} \rangle$	\hat{s}_{P+3}
$\langle s_P, \hat{s}_{P+1}, \hat{s}_{P+2}, \hat{s}_{P+3} \rangle$	\hat{s}_{P+4}
$\langle \hat{s}_{P+1}, \hat{s}_{P+2}, \hat{s}_{P+3}, \hat{s}_{P+4} \rangle$	\hat{s}_{P+5}
\vdots	\vdots

Ingesting sequential I/O data

Text generation

- Sequence of P characters: $\langle s_0, s_1, s_2, \dots, s_P \rangle$
- s_p : the p th character

Text generation

- Sequence of P characters: $\langle s_0, s_1, s_2, \dots, s_P \rangle$
- s_p : the p th character
- - **classification problem**- I/O are windowed subsequences

Text generation


- Sequence of P characters: $\langle s_0, s_1, s_2, \dots, s_P \rangle$
- s_p : the p th character
- **classification problem**- I/O are windowed subsequences
 - so need training and testing sets
 - lets see how both are formed, start with training

d o g s a r e g r e a t


dogs are great

input output


d o g s a r e g r e a t




d o g s a r e g r e a t




dogs are great




d o g s a r e g r e a t




d o g s a r e g r e a t




d o g s a r e g r e a t



d o g s a r e g r e a t



dogs are great



d o g s a r e g r e a t

←————→ ←→
input output

training

testing

One more time - with notation

dogs are great

← input → output →

Input

$\langle s_0, s_1, s_2, s_3 \rangle$

Output

s_4

d o g s a r e g r e a t

←————→ ↔
input output

Input	Output
$\langle s_0, s_1, s_2, s_3 \rangle$	s_4
$\langle s_1, s_2, s_3, s_4 \rangle$	s_5

d o g s a r e g r e a t

← input → output →

Input	Output
$\langle s_0, s_1, s_2, s_3 \rangle$	s_4
$\langle s_1, s_2, s_3, s_4 \rangle$	s_5
\vdots	\vdots

dogs are great

← input → output →


Input	Output
$\langle s_0, s_1, s_2, s_3 \rangle$	s_4
$\langle s_1, s_2, s_3, s_4 \rangle$	s_5
\vdots	\vdots

d o g s **a r e** **g** r e a t

←————→ ↔
input output


Input	Output
$\langle s_0, s_1, s_2, s_3 \rangle$	s_4
$\langle s_1, s_2, s_3, s_4 \rangle$	s_5
\vdots	\vdots

dogs are great




Input	Output
$\langle s_0, s_1, s_2, s_3 \rangle$	s_4
$\langle s_1, s_2, s_3, s_4 \rangle$	s_5
\vdots	\vdots

d o g s a r e g r e a t




Input	Output
$\langle s_0, s_1, s_2, s_3 \rangle$	s_4
$\langle s_1, s_2, s_3, s_4 \rangle$	s_5
\vdots	\vdots

dogs are great




Input	Output
$\langle s_0, s_1, s_2, s_3 \rangle$	s_4
$\langle s_1, s_2, s_3, s_4 \rangle$	s_5
\vdots	\vdots

dogs are great



Input	Output
$\langle s_0, s_1, s_2, s_3 \rangle$	s_4
$\langle s_1, s_2, s_3, s_4 \rangle$	s_5
\vdots	\vdots
$\langle s_{P-4}, s_{P-3}, s_{P-2}, s_{P-1} \rangle$	s_P

dogs are great



training

testing

dogs are great



input

training

testing

dogs are **greate**

←→
input prediction

Input	Output
$\langle s_{P-3}, s_{P-2}, s_{P-1}, s_P \rangle$	\hat{s}_{P+1}

d o g s a r e g r e a t e



input


Input	Output
$\langle s_{P-3}, s_{P-2}, s_{P-1}, s_P \rangle$	\hat{s}_{P+1}

dogs are greater

←→
input prediction

Input	Output
$\langle s_{P-3}, s_{P-2}, s_{P-1}, s_P \rangle$	\hat{s}_{P+1}
$\langle s_{P-2}, s_{P-1}, s_P, \hat{s}_{P+1} \rangle$	\hat{s}_{P+2}

d o g s a r e g r e a t e r




Input	Output
$\langle s_{P-3}, s_{P-2}, s_{P-1}, s_P \rangle$	\hat{s}_{P+1}
$\langle s_{P-2}, s_{P-1}, s_P, \hat{s}_{P+1} \rangle$	\hat{s}_{P+2}

dogs are greater

input prediction


Input	Output
$\langle s_{P-3}, s_{P-2}, s_{P-1}, s_P \rangle$	\hat{s}_{P+1}
$\langle s_{P-2}, s_{P-1}, s_P, \hat{s}_{P+1} \rangle$	\hat{s}_{P+2}
$\langle s_{P-1}, s_P, \hat{s}_{P+1}, \hat{s}_{P+2} \rangle$	\hat{s}_{P+3}

dogs are greater




Input	Output
$\langle s_{P-3}, s_{P-2}, s_{P-1}, s_P \rangle$	\hat{s}_{P+1}
$\langle s_{P-2}, s_{P-1}, s_P, \hat{s}_{P+1} \rangle$	\hat{s}_{P+2}
$\langle s_{P-1}, s_P, \hat{s}_{P+1}, \hat{s}_{P+2} \rangle$	\hat{s}_{P+3}

d o g s a r e g r e a t e r t



Input	Output
$\langle s_{P-3}, s_{P-2}, s_{P-1}, s_P \rangle$	\hat{s}_{P+1}
$\langle s_{P-2}, s_{P-1}, s_P, \hat{s}_{P+1} \rangle$	\hat{s}_{P+2}
$\langle s_{P-1}, s_P, \hat{s}_{P+1}, \hat{s}_{P+2} \rangle$	\hat{s}_{P+3}
$\langle s_P, \hat{s}_{P+1}, \hat{s}_{P+2}, \hat{s}_{P+3} \rangle$	\hat{s}_{P+4}

d o g s a r e g r e a t e r t



Input	Output
$\langle s_{P-3}, s_{P-2}, s_{P-1}, s_P \rangle$	\hat{s}_{P+1}
$\langle s_{P-2}, s_{P-1}, s_P, \hat{s}_{P+1} \rangle$	\hat{s}_{P+2}
$\langle s_{P-1}, s_P, \hat{s}_{P+1}, \hat{s}_{P+2} \rangle$	\hat{s}_{P+3}
$\langle s_P, \hat{s}_{P+1}, \hat{s}_{P+2}, \hat{s}_{P+3} \rangle$	\hat{s}_{P+4}

dogs are greater than



input prediction


Input	Output
$\langle s_{P-3}, s_{P-2}, s_{P-1}, s_P \rangle$	\hat{s}_{P+1}
$\langle s_{P-2}, s_{P-1}, s_P, \hat{s}_{P+1} \rangle$	\hat{s}_{P+2}
$\langle s_{P-1}, s_P, \hat{s}_{P+1}, \hat{s}_{P+2} \rangle$	\hat{s}_{P+3}
$\langle s_P, \hat{s}_{P+1}, \hat{s}_{P+2}, \hat{s}_{P+3} \rangle$	\hat{s}_{P+4}
$\langle \hat{s}_{P+1}, \hat{s}_{P+2}, \hat{s}_{P+3}, \hat{s}_{P+4} \rangle$	\hat{s}_{P+5}

dogs are greater th

input

Input	Output
$\langle s_{P-3}, s_{P-2}, s_{P-1}, s_P \rangle$	\hat{s}_{P+1}
$\langle s_{P-2}, s_{P-1}, s_P, \hat{s}_{P+1} \rangle$	\hat{s}_{P+2}
$\langle s_{P-1}, s_P, \hat{s}_{P+1}, \hat{s}_{P+2} \rangle$	\hat{s}_{P+3}
$\langle s_P, \hat{s}_{P+1}, \hat{s}_{P+2}, \hat{s}_{P+3} \rangle$	\hat{s}_{P+4}
$\langle \hat{s}_{P+1}, \hat{s}_{P+2}, \hat{s}_{P+3}, \hat{s}_{P+4} \rangle$	\hat{s}_{P+5}

dogs are greater than


input prediction

Input	Output
$\langle s_{P-3}, s_{P-2}, s_{P-1}, s_P \rangle$	\hat{s}_{P+1}
$\langle s_{P-2}, s_{P-1}, s_P, \hat{s}_{P+1} \rangle$	\hat{s}_{P+2}
$\langle s_{P-1}, s_P, \hat{s}_{P+1}, \hat{s}_{P+2} \rangle$	\hat{s}_{P+3}
$\langle s_P, \hat{s}_{P+1}, \hat{s}_{P+2}, \hat{s}_{P+3} \rangle$	\hat{s}_{P+4}
$\langle \hat{s}_{P+1}, \hat{s}_{P+2}, \hat{s}_{P+3}, \hat{s}_{P+4} \rangle$	\hat{s}_{P+5}
\vdots	\vdots

here we illustrated with
window size $T = 4$

d o g s a r e g r e a t e r t h a n



Input	Output
$\langle s_{P-3}, s_{P-2}, s_{P-1}, s_P \rangle$	\hat{s}_{P+1}
$\langle s_{P-2}, s_{P-1}, s_P, \hat{s}_{P+1} \rangle$	\hat{s}_{P+2}
$\langle s_{P-1}, s_P, \hat{s}_{P+1}, \hat{s}_{P+2} \rangle$	\hat{s}_{P+3}
$\langle s_P, \hat{s}_{P+1}, \hat{s}_{P+2}, \hat{s}_{P+3} \rangle$	\hat{s}_{P+4}
$\langle \hat{s}_{P+1}, \hat{s}_{P+2}, \hat{s}_{P+3}, \hat{s}_{P+4} \rangle$	\hat{s}_{P+5}
\vdots	\vdots

d o g s a r e g r e a t e r t h a n

(character pre-processing)


- characters \rightarrow numbers for supervised models
- use one-hot encoding scheme


$$a \leftarrow \begin{bmatrix} 1 \\ 0 \\ 0 \\ \vdots \\ 0 \\ 0 \end{bmatrix} \quad b \leftarrow \begin{bmatrix} 0 \\ 1 \\ 0 \\ \vdots \\ 0 \\ 0 \end{bmatrix} \quad c \leftarrow \begin{bmatrix} 0 \\ 0 \\ 1 \\ \vdots \\ 0 \\ 0 \end{bmatrix} \dots$$

RNN architecture


RNN architecture

- **Want:** parameterized supervised learning model that enforces ordered sequential-ness on I/O
- So: extend feedforward nets to ingest sequences





DNN



RNN

NN

feed input in
as vector

$$h = \tanh \left(v_0 + \sum_{t=1}^T v_t s_t \right)$$

$$h_0 = \tanh (v_0 + v s_0)$$

$$h_t = \tanh (v_0 + v s_t + u h_{t-1})$$

$$\hat{y} = b + w h$$

$$\hat{y} = b + w h_T$$

$$(y - \hat{y})^2$$

$$(y - \hat{y})^2$$

RNN technical issues

- RNN still trained via gradient descent (a.k.a. backprop)
- Similar 'vanishing gradient' problem as with feedforward nets
- Using different activation (relu) helps, but additional unit architecture helps a lot (Long Term Short Memory module)

SUMMARY

Supervised learning

+

structured data

Beyond vanilla

- Vanilla models don't exploit ordered sequential I/O
- Include I/O structure in learning framework —> better results
 - Engineer into fixed feature extraction
 - limbed parameterized feature extractor in model (e.g., convnets, RNNs)
- RNNs one parameterized way to exploit sequential I/O
- RNNs are natural extension of feedforward nets, and inherit similar technical issues