

Bayesian Optimization with scikit-optimize

Tim Head

31 August 2017



How to Make The Best



Beer recipes

Your task: brew the most tasty beer possible.

Many parameters you can tweak, for simplicity we let's pretend there are only two: **alcohol content** and **bitterness**.

How do you find the best combination?



Evaluating a recipe is expensive

How to score a beer:


- Buy ingredients
- Brew it
- ...wait...
- Find expert panel and collect scores

This means we can't try a large number of combinations, have to be smart.

This is an optimization problem, with a (very) expensive *objective function*.

scipy.optimize!

Bayesian Optimisation



Demo

$$x^* = \arg \max_x f(x)$$

- f is a black box function, with no closed form nor gradients.
- f is expensive to evaluate.
- You only have noisy observations of f .

If you do not have
these constraints,
do not use
Bayesian
optimization.

Back to beer

scikit-optimize



Andreas Mueller

@amuellerm1

Following



Trying to review Bayesian optimization packages for Python. So far: 3/6 installed according to instructions, 0/6 passed tests on my machine.

9:49 PM - 2 Aug 2017

10 Retweets 73 Likes



```
$ pip install numpy
```

```
$ pip install scikit-optimize
```

Like scipy.optimize

```
from skopt import gp_minimize  
  
res = gp_minimize(f, [(-2.0, 2.0)])
```

Ask-and-tell interface

```
from skopt import Optimizer

opt = Optimizer([(-2.0, 2.0)])
# get a new suggestion
suggested = opt.ask()
# evaluate the suggestion
y = f(suggested)
# give feedback to the optimizer
opt.tell(suggested, y)
```


With scikit-learn


```
from skopt import BayesSearchCV  
  
bayes = BayesSearchCV(clf, n_iter=32)  
bayes.fit(X_train, y_train)  
print(bayes.cv_results_)
```

```
bounds = [(-5.0, 10.0), (0.0, 15.0)]
```

```
n_calls = 160
```

```
forest_res = forest_minimize(branin, bounds, n_calls=n_calls,
                             base_estimator="ET",
                             random_state=4)
```

```
_ = plot_evaluations(forest_res, bins=10)
```



1. Given observations $(x_i, y_i = f(x_i))$ for $i = 1, \dots, t$, integrate out all possible true function f .
2. Optimize a cheap acquisition/utility function to select the next point.

Exploit uncertainty to balance exploration

3. Sample the next observation y_{t+1} at x_{t+1} .

Acquisition functions

Acquisition functions $u(x)$ specify which sample

- Expected improvement (default): $-EI(x)$
- Lower confidence bound: $LCB(x) = \mu_{GP}(x) - \sigma_{GP}(x)$
- Probability of improvement: $-PI(x) = 1 - \Phi(\frac{\mu_{GP}(x) - y_{t+1}^+}{\sigma_{GP}(x)})$

where x_t^+ is the best point observed so far.

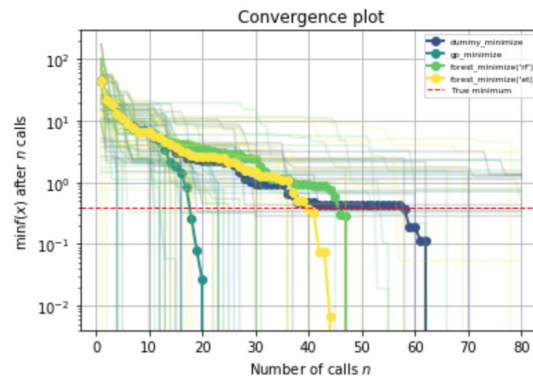
In most cases, acquisition functions provide knobs (e.g., κ) for controlling the exploration-exploitation trade-off. - Search in regions where $\mu_{GP}(x)$ is high (exploitation) - Probe regions where uncertainty $\sigma_{GP}(x)$ is high (exploration)

Note that this can take a few minutes.

```
from skopt.plots import plot_convergence
```

```
plot = plot_convergence(("dummy_minimize", dummy_res),
                        ("gp_minimize", gp_res),
                        ("forest_minimize('rf')", rf_res),
                        ("forest_minimize('et')", et_res),
                        true_minimum=0.397887, ylabel="log")
```

```
plot.legend(loc="best", prop={'size': 6}, numpoints=1);
```



Join us (not just tomorrow)

 scikit-optimize / scikit-optimize

 Unwatch ▾

43

★ Unstar

562



<> Code

⚠ Issues 78

🔗 Pull requests 11

📁 Projects 0

📖 Wiki

⚙ Settings

🔍 Insights ▾

Sequential model-based optimization with a `scipy.optimize` interface <https://scikit-optimize.github.io>

bayesopt

optimization

scientific-computing

python

machine-learning

hyperparameter

bayesian-optimization

Manage topics

📦 802 commits

🌿 1 branch

📦 5 releases

👤 23 contributors

??!?

Tim Head

✉ tim@wildtreotech.com

🐦 @betatim

To run the notebooks:

<https://github.com/wildtreotech/bayesian-optimisation>

Brewing beer is
expensive and
does not come
with gradients,
scikit-optimize can help.