



Matrix Multiplication

Chapter I – Matrix Multiplication

By Gokturk Poyrazoglu

The State University of New York at Buffalo – BEST Group – Winter Lecture Series

Outline

1. Basic Algorithms and Notation
2. Structure and Efficiency
3. Block Matrices
4. Matrix-Vector Products
5. Parallel Matrix Multiplication



Matrix Notation

$$A \in \mathbb{R}^{m \times n} \iff A = (a_{ij}) = \begin{bmatrix} a_{11} & \cdots & a_{1n} \\ \vdots & & \vdots \\ a_{m1} & \cdots & a_{mn} \end{bmatrix}, \quad a_{ij} \in \mathbb{R}$$

▶ Matrix Operations

- ▶ Transposition : $C = A^T \implies c_{ij} = a_{ji},$
- ▶ Addition : $C = A + B \implies c_{ij} = a_{ij} + b_{ij},$
- ▶ Scalar-matrix Multiplication : $C = \alpha A \implies c_{ij} = \alpha a_{ij},$
- ▶ Matrix-matrix Multiplication : $C = AB \implies c_{ij} = \sum_{k=1}^p a_{ik} b_{kj}.$
- ▶ Pointwise Multiplication : $C = A .* B \implies c_{ij} = a_{ij} b_{ij}$
- ▶ Pointwise Division : $C = A ./ B \implies c_{ij} = a_{ij} / b_{ij}.$



Vector Notation

► Column Vector :

$$x \in \mathbb{R}^n \iff x = \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix} \quad x_i \in \mathbb{R}.$$

► Row Vector :

$$x \in \mathbb{R}^{1 \times n} \iff x = [x_1, \dots, x_n].$$



Vector Operations

- ▶ Scalar-vector Multiplication :

$$z = ax \implies z_i = ax_i,$$

- ▶ Vector Addition :

$$z = x + y \implies z_i = x_i + y_i,$$

- ▶ Inner Product (dot product) :

$$c = x^T y \implies c = \sum_{i=1}^n x_i y_i.$$

- ▶ Vector Update :

$$y = ax + y \implies y_i = ax_i + y_i$$

- ▶ Pointwise Vector Multiplication :

$$z = x .* y \implies z_i = x_i y_i,$$

- ▶ Pointwise Vector Division :

$$z = x ./ y \implies z_i = x_i / y_i.$$



Dot Product Algorithm

- ▶ Dot product:

$$c = x^T y \quad \Rightarrow \quad c = \sum_{i=1}^n x_i y_i.$$

- ▶ Algorithm :

```
c = 0
for i = 1:n
    c = c + x(i)y(i)
end
```



n multiplications

n additions

- ▶ The dot product operation is an $O(n)$ operation.
- ▶ *The amount of work scales linearly with the dimension*



Saxpy (Vector update) Algorithm

► Vector update : $y = ax + y \implies y_i = ax_i + y_i$

► Algorithm :

```
for i = 1:n  
    y(i) = y(i) + ax(i)  
end
```

- The vector update operation is also an $O(n)$ operation.
- *The amount of work scales linearly with the dimension*



Matrix-vector Multiplication

$$y = y + Ax \quad \longrightarrow \quad y_i = y_i + \sum_{j=1}^n a_{ij}x_j, \quad i = 1:m,$$

where y and x are vectors and A is a matrix

► Algorithm:

```
for  $i = 1:m$ 
    for  $j = 1:n$ 
         $y(i) = y(i) + A(i,j)x(j)$ 
    end
end
```

- This algorithm is called Row-oriented gaxpy.
 - This algorithm involves $O(nm)$ work.
-



Matrix-vector Multiplication

$$y = y + Ax$$

where y and x are vectors and A is a matrix.

► 2nd Approach : Column Oriented Gaxpy

Example :

$$\begin{bmatrix} 1 & 2 \\ 3 & 4 \\ 5 & 6 \end{bmatrix} \begin{bmatrix} 7 \\ 8 \end{bmatrix} = \begin{bmatrix} 1 \cdot 7 + 2 \cdot 8 \\ 3 \cdot 7 + 4 \cdot 8 \\ 5 \cdot 7 + 6 \cdot 8 \end{bmatrix} = 7 \begin{bmatrix} 1 \\ 3 \\ 5 \end{bmatrix} + 8 \begin{bmatrix} 2 \\ 4 \\ 6 \end{bmatrix} = \begin{bmatrix} 23 \\ 53 \\ 83 \end{bmatrix}$$

Algorithm :

```
for j = 1:n  
    for i = 1:m  
        y(i) = y(i) + A(i,j) * x(j)  
    end  
end
```

Interchanging the
order of the loop



Partitioning a Matrix

- ▶ Consider a matrix A as a stack of row vectors

$$A \in \mathbb{R}^{m \times n} \quad \Longleftrightarrow \quad A = \begin{bmatrix} r_1^T \\ \vdots \\ r_m^T \end{bmatrix}, \quad r_k \in \mathbb{R}^n.$$

- ▶ Row Partition Example:

- ▶ $A = \begin{bmatrix} 1 & 2 \\ 3 & 4 \\ 5 & 6 \end{bmatrix}$, so that $r_1^T = [1 \ 2]$, $r_2^T = [3 \ 4]$, $r_3^T = [5 \ 6]$.

- ▶ New algorithm for vector update

```
for  $i = 1:m$   
     $y_i = y_i + r_i^T x$   
end
```



Partitioning a Matrix

- ▶ Column Partition Example

- ▶ $A = \begin{bmatrix} 1 & 2 \\ 3 & 4 \\ 5 & 6 \end{bmatrix}$, so that $c_1 = \begin{bmatrix} 1 \\ 3 \\ 5 \end{bmatrix}$, $c_2 = \begin{bmatrix} 2 \\ 4 \\ 6 \end{bmatrix}$

- ▶ Algorithm for vector update with column partition

```
for  $j = 1:n$ 
```

```
     $y = y + x_j c_j$ 
```

```
end
```

Count number of columns

Multiply a scalar with a vector

Colon Notation

- ▶ k^{th} row of matrix A:

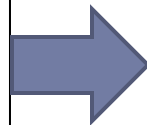
$$A(k, :) = [a_{k1}, \dots, a_{kn}] .$$

- ▶ k^{th} column of matrix A:

$$A(:, k) = \begin{bmatrix} a_{1k} \\ \vdots \\ a_{mk} \end{bmatrix} .$$

- ▶ Rewrite the algorithm with row partitioning

```
for  $i = 1:m$   
     $y_i = y_i + \underline{r_i^T} x$   
end
```



```
for  $i = 1:m$   
     $y(i) = y(i) + \underline{A(i, :)} \cdot x$   
end
```



Outer Product Update

- ▶ Update matrix A with multiplication of two vectors

$$A = A + xy^T, \quad A \in \mathbb{R}^{m \times n}, x \in \mathbb{R}^m, y \in \mathbb{R}^n.$$

- ▶ Algorithm for outer product update:

```
for  $i = 1:m$ 
  for  $j = 1:n$ 
     $a_{ij} = a_{ij} + x_i y_j$ 
  end
end
```

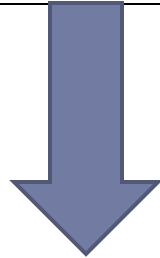
- ▶ The algorithm involves $O(nm)$ operations.



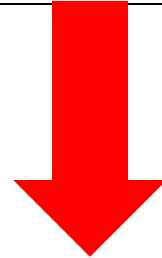
Outer Product Update with Vector Mult.

```
for  $i = 1:m$   
  for  $j = 1:n$   
     $a_{ij} = a_{ij} + x_i y_j$   
  end  
end
```

Row update



```
for  $i = 1:m$   
   $A(i, :) = A(i, :) + x(i) \cdot y^T$   
end
```



Column Update

```
for  $j = 1:n$   
   $A(:, j) = A(:, j) + y(j) \cdot x$   
end
```



Matrix – Matrix Multiplication

► Dot Product

$$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} \begin{bmatrix} 5 & 6 \\ 7 & 8 \end{bmatrix} = \begin{bmatrix} 1 \cdot 5 + 2 \cdot 7 & 1 \cdot 6 + 2 \cdot 8 \\ 3 \cdot 5 + 4 \cdot 7 & 3 \cdot 6 + 4 \cdot 8 \end{bmatrix}$$

► Linear Combination of left matrix columns

$$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} \begin{bmatrix} 5 & 6 \\ 7 & 8 \end{bmatrix} = \begin{bmatrix} 5 \begin{bmatrix} 1 \\ 3 \end{bmatrix} + 7 \begin{bmatrix} 2 \\ 4 \end{bmatrix}, & 6 \begin{bmatrix} 1 \\ 3 \end{bmatrix} + 8 \begin{bmatrix} 2 \\ 4 \end{bmatrix} \end{bmatrix}.$$

► Sum of Outer Products

$$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} \begin{bmatrix} 5 & 6 \\ 7 & 8 \end{bmatrix} = \begin{bmatrix} 1 \\ 3 \end{bmatrix} \begin{bmatrix} 5 & 6 \end{bmatrix} + \begin{bmatrix} 2 \\ 4 \end{bmatrix} \begin{bmatrix} 7 & 8 \end{bmatrix}$$



Matrix – Matrix Multiplication

- ▶ Consider a matrix update by matrix multiplication

$$C = C + AB, \quad C \in \mathbb{R}^{m \times n}, A \in \mathbb{R}^{m \times r}, B \in \mathbb{R}^{r \times n}.$$

- ▶ Triply Nested Loop Algorithm


```
for i = 1:m
    for j = 1:n
        for k = 1:r
            C(i,j) = C(i,j) + A(i,k)·B(k,j)
        end
    end
end
```

- ▶ This algorithm involves $O(mnr)$ operations.
-

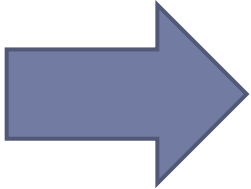


Dot Product Matrix Multiplication


```
for i = 1:m
    for j = 1:n
         $C(i, j) = C(i, j) + A(i, :) \cdot B(:, j)$ 
    end
end
```



```
for i = 1:m
    for j = 1:n
         $c_{ij} = c_{ij} + a_i^T b_j$ 
    end
end
```



```
for i = 1:m
     $c_i^T = c_i^T + a_i^T B$ 
end
```



```
for i = 1:m
     $C(i, :) = C(i, :) + A(i, :) \cdot B$ 
end
```



Saxpy Formulation

- ▶ Suppose matrix-A and C are column partitioned

$$A = [a_1 | \cdots | a_r], \quad C = [c_1 | \cdots | c_n].$$

- ▶ Each column of C can be updated as follows :

$$c_j = c_j + \sum_{k=1}^r a_k b_{kj}, \quad j = 1:n.$$

Algorithm :

```
for j = 1:n
    for k = 1:r
        C(:,j) = C(:,j) + A(:,k)·B(k,j)
    end
end
```

Simplified Algorithm:

```
for j = 1:n
    C(:,j) = C(:,j) + AB(:,j)
end
```



Flops

Flop is a measure of;

1. Addition
2. Subtraction
3. Multiplication
4. Division

Operation	Dimension	Flops
$\alpha = x^T y$	$x, y \in \mathbb{R}^n$	$2n$
$y = y + ax$	$a \in \mathbb{R}, x, y \in \mathbb{R}^n$	$2n$
$y = y + Ax$	$A \in \mathbb{R}^{m \times n}, x \in \mathbb{R}^n, y \in \mathbb{R}^m$	$2mn$
$A = A + yx^T$	$A \in \mathbb{R}^{m \times n}, x \in \mathbb{R}^n, y \in \mathbb{R}^m$	$2mn$
$C = C + AB$	$A \in \mathbb{R}^{m \times r}, B \in \mathbb{R}^{r \times n}, C \in \mathbb{R}^{m \times n}$	$2mnr$



Complex Matrices

- ▶ Consider matrix A :

$$A = B + iC \in \mathbb{C}^{m \times n},$$

- ▶ B is the real part of A ;
- ▶ C is the imaginary part of A
- ▶ Matrix Operations:
 - ▶ Transposition of A becomes *conjugate transposition*.

$$C = A^H \implies c_{ij} = \bar{a}_{ji}.$$



Structure and Efficiency

Outline

1. Band Matrices
2. Triangular Matrices
3. Diagonal Matrices
4. Symmetric Matrices
5. Permutation Matrices



Band Matrices

- ▶ Consider a matrix A ;

$$A = \begin{bmatrix} \times & \times & \times & 0 & 0 \\ \times & \times & \times & \times & 0 \\ 0 & \times & \times & \times & \times \\ 0 & 0 & \times & \times & \times \\ 0 & 0 & 0 & \times & \times \\ 0 & 0 & 0 & 0 & \times \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}.$$

- ▶ Matrix A has *lower bandwidth* p if $a_{ij}=0$ when $i > j + p$
- ▶ Matrix A has *upper bandwidth* q if $a_{ij}=0$ when $j > i + q$



Triangular Matrix Multiplication

- ▶ Consider a matrix update by matrix multiplication

$$AB = \begin{bmatrix} a_{11}b_{11} & a_{11}b_{12} + a_{12}b_{22} & a_{11}b_{13} + a_{12}b_{23} + a_{13}b_{33} \\ 0 & a_{22}b_{22} & a_{22}b_{23} + a_{23}b_{33} \\ 0 & 0 & a_{33}b_{33} \end{bmatrix}.$$

- ▶ So the update takes the form of

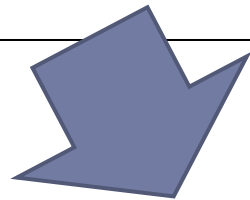
$$c_{ij} = c_{ij} + \sum_{k=i}^j a_{ik}b_{kj}$$

- ▶ k starts from i to j so we ignore the zero elements of [AB]



Colon Notation Triangular Matrix Mult.

```
for i = 1:n
    for j = i:n
        for k = i:j
             $C(i, j) = C(i, j) + A(i, k) \cdot B(k, j)$ 
        end
    end
end
```



```
for i = 1:n
    for j = i:n
         $C(i, j) = C(i, j) + A(i, i:j) \cdot B(i:j, j)$ 
    end
end
```


Diagonal Matrices

- ▶ Lower and upper bandwidth is Zero for all diagonal matrices.
- ▶ Notation :

$$D = \text{diag}(d_1, \dots, d_q), \quad q = \min\{m, n\} \quad \Longleftrightarrow \quad d_i = d_{ii}.$$

- ▶ Pre-multiplication of D scales rows of A

$$B = DA \quad \Longleftrightarrow \quad B(i, :) = d_i \cdot A(i, :), \quad i = 1:m$$

- ▶ Post-multiplication of D scales columns of A

$$B = AD \quad \Longleftrightarrow \quad B(:, j) = d_j \cdot A(:, j), \quad j = 1:n.$$



Symmetry

$$A^T = A$$

Symmetric: $\begin{bmatrix} 1 & 2 & 3 \\ 2 & 4 & 5 \\ 3 & 5 & 6 \end{bmatrix},$

$$A^H = A$$

Hermitian: $\begin{bmatrix} 1 & 2-3i & 4-5i \\ 2+3i & 6 & 7-8i \\ 4+5i & 7+8i & 9 \end{bmatrix},$

$$A^T = -A$$

Skew-Symmetric: $\begin{bmatrix} 0 & -2 & 3 \\ 2 & 0 & -5 \\ -3 & 5 & 0 \end{bmatrix},$

$$A^H = -A$$

Skew-Hermitian: $\begin{bmatrix} i & -2+3i & -4+5i \\ 2+3i & 6i & -7+8i \\ 4+5i & 7+8i & 9i \end{bmatrix}$



Identity and Permutation Matrix

- ▶ Identity Matrix:

$$I_4 = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}.$$

- ▶ Permutation Matrix

$$P = \begin{bmatrix} 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 \\ 1 & 0 & 0 & 0 \end{bmatrix}$$



Famous Permutation Matrices

► Exchange Permutation

$$y = \mathcal{E}_4 x = \begin{bmatrix} 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 \\ 0 & 1 & 0 & 0 \\ 1 & 0 & 0 & 0 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{bmatrix} = \begin{bmatrix} x_4 \\ x_3 \\ x_2 \\ x_1 \end{bmatrix}$$

► Downshift Permutation

$$y = \mathcal{D}_4 x = \begin{bmatrix} 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{bmatrix} = \begin{bmatrix} x_4 \\ x_1 \\ x_2 \\ x_3 \end{bmatrix}$$



Block Matrix Terminology

- ▶ Special case of column and row partitioning.

$$A = \left[\begin{array}{ccc} A_{11} & \dots & A_{1r} \\ \vdots & & \vdots \\ A_{q1} & \dots & A_{qr} \end{array} \right] \begin{array}{l} m_1 \\ \\ m_q \end{array}$$

$n_1 \qquad n_r$

- ▶ Examples

$$\text{diag}(A_{11}, A_{22}, A_{33}) = \left[\begin{array}{ccc} A_{11} & 0 & 0 \\ 0 & A_{22} & 0 \\ 0 & 0 & A_{33} \end{array} \right]$$

$$L = \left[\begin{array}{ccc} L_{11} & 0 & 0 \\ L_{21} & L_{22} & 0 \\ L_{31} & L_{32} & L_{33} \end{array} \right], \quad U = \left[\begin{array}{ccc} U_{11} & U_{12} & U_{13} \\ 0 & U_{22} & U_{23} \\ 0 & 0 & U_{33} \end{array} \right], \quad T = \left[\begin{array}{ccc} T_{11} & T_{12} & 0 \\ T_{21} & T_{22} & T_{23} \\ 0 & T_{32} & T_{33} \end{array} \right]$$



Block Matrix Operations

- Scalar Multiplication :

$$\mu \begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \\ A_{31} & A_{32} \end{bmatrix} = \begin{bmatrix} \mu A_{11} & \mu A_{12} \\ \mu A_{21} & \mu A_{22} \\ \mu A_{31} & \mu A_{32} \end{bmatrix}$$

- Transposition :

$$\begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \\ A_{31} & A_{32} \end{bmatrix}^T = \begin{bmatrix} A_{11}^T & A_{21}^T & A_{31}^T \\ A_{12}^T & A_{22}^T & A_{32}^T \end{bmatrix}$$

- Addition:

$$\begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \\ A_{31} & A_{32} \end{bmatrix} + \begin{bmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \\ B_{31} & B_{32} \end{bmatrix} = \begin{bmatrix} A_{11} + B_{11} & A_{12} + B_{12} \\ A_{21} + B_{21} & A_{22} + B_{22} \\ A_{31} + B_{31} & A_{32} + B_{32} \end{bmatrix}$$



Block Matrix Multiplication

$$\begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \\ A_{31} & A_{32} \end{bmatrix} \begin{bmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{bmatrix} = \begin{bmatrix} A_{11}B_{11} + A_{12}B_{21} & A_{11}B_{12} + A_{12}B_{22} \\ A_{21}B_{11} + A_{22}B_{21} & A_{21}B_{12} + A_{22}B_{22} \\ A_{31}B_{11} + A_{32}B_{21} & A_{31}B_{12} + A_{32}B_{22} \end{bmatrix}$$

- ▶ Column dimensions of A
must match with
row dimensions of B.



Block Vector Update

- ▶ Consider Block Matrix A and a block vector y

$$A = \begin{bmatrix} A_1 \\ \vdots \\ A_q \end{bmatrix} \begin{matrix} m_1 \\ \\ m_q \end{matrix}$$

$$y = \begin{bmatrix} y_1 \\ \vdots \\ y_q \end{bmatrix} \begin{matrix} m_1 \\ \\ m_q \end{matrix}$$

- ▶ **Block Vector update :**
- ▶ **Algorithm :**

```
 $\alpha = 0$   
for  $i = 1:q$   
     $idx = \alpha + 1 : \alpha + m_i$   
     $y(idx) = y(idx) + A(idx, :) \cdot x$   
     $\alpha = \alpha + m_i$   
end
```

$$\begin{bmatrix} y_1 \\ \vdots \\ y_q \end{bmatrix} = \begin{bmatrix} y_1 \\ \vdots \\ y_q \end{bmatrix} + \begin{bmatrix} A_1 \\ \vdots \\ A_q \end{bmatrix} x$$

Block Matrix Update

- Consider a matrix update by matrix multiplication.

$$C = C + AB$$

$A = (A_{\alpha\beta})$, $B = (B_{\alpha\beta})$, and $C = (C_{\alpha\beta})$ as N -by- N block matrices with ℓ -by- ℓ blocks.

$$C_{\alpha\beta} = C_{\alpha\beta} + \sum_{\gamma=1}^N A_{\alpha\gamma} B_{\gamma\beta}, \quad \alpha = 1:N, \quad \beta = 1:N.$$

- Algorithm:

```
for  $\alpha = 1:N$ 
     $i = (\alpha - 1)\ell + 1:\alpha\ell$ 
    for  $\beta = 1:N$ 
         $j = (\beta - 1)\ell + 1:\beta\ell$ 
        for  $\gamma = 1:N$ 
             $k = (\gamma - 1)\ell + 1:\gamma\ell$ 
             $C(i, j) = C(i, j) + A(i, k) \cdot B(k, j)$ 
        end
    end
end
```

Complex Matrix Multiplication

- ▶ Consider matrix A, B, and C are complex matrices where all matrices are real and $i^2 = -1$

$$C_1 + iC_2 = (C_1 + iC_2) + (A_1 + iA_2)(B_1 + iB_2)$$

- ▶ Multiplication:

$$\begin{bmatrix} C_1 \\ C_2 \end{bmatrix} = \begin{bmatrix} C_1 \\ C_2 \end{bmatrix} + \begin{bmatrix} A_1 & -A_2 \\ A_2 & A_1 \end{bmatrix} \begin{bmatrix} B_1 \\ B_2 \end{bmatrix}$$

- ▶ Note: Complex Matrix Multiplication has expanded dimension.



Hamiltonian Matrix

► Form :

$$M = \begin{bmatrix} A & G \\ F & -A^T \end{bmatrix}$$

1. A, F, and G are square matrices
2. F and G are symmetric matrices

Hamiltonian Check

1. Consider a permutation matrix J as

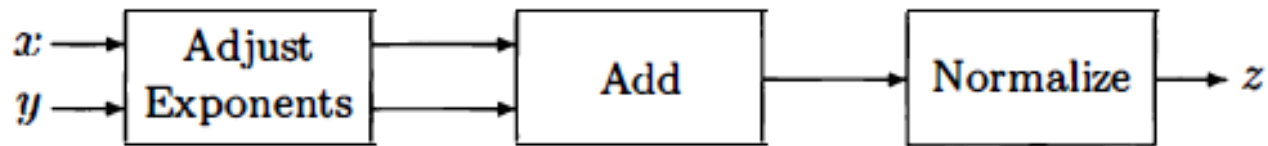
$$J = \begin{bmatrix} 0 & I_n \\ -I_n & 0 \end{bmatrix}$$

2. If $JMJ^T = -M^T$, then M is Hamiltonian Matrix.

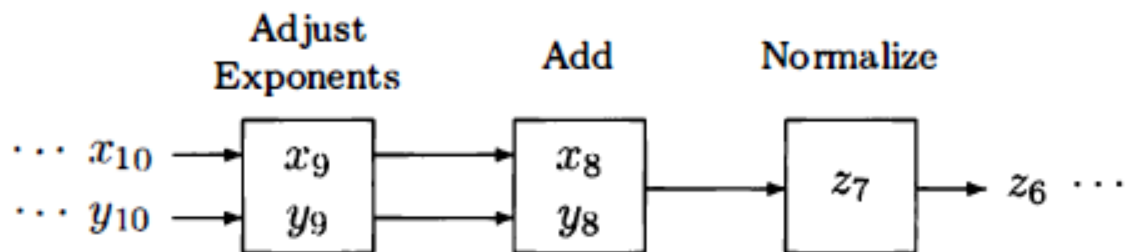


Vector Processing

► 3-cycle Adder



► Pipelined Addition for a vector operation



Data Motion in Computation Performance

Other than flop counts;

Data motion is also an important factor
when reasoning about performance.

```
first = 1
while first ≤ n
    last = min{n, first + vL − 1}
    Vector load:  $r_1 \leftarrow x(\textit{first}:\textit{last})$ 
    Vector load:  $r_2 \leftarrow y(\textit{first}:\textit{last})$ 
    Vector add:  $r_1 = r_1 + r_2$ 
    Vector store:  $z(\textit{first}:\textit{last}) \leftarrow r_1$ 
    first = last + 1
end
```



Load/Store Operations

Vector Update (Gaxpy)

$$y = y + Ax$$

Load / Store Operations: (3+n)

```
 $r_x \leftarrow x$   
 $r_y \leftarrow y$   
for  $j = 1:n$   
     $r_a \leftarrow A(:, j)$   
     $r_y = r_y + r_a r_x(j)$   
end  
 $y \leftarrow r_y$ 
```

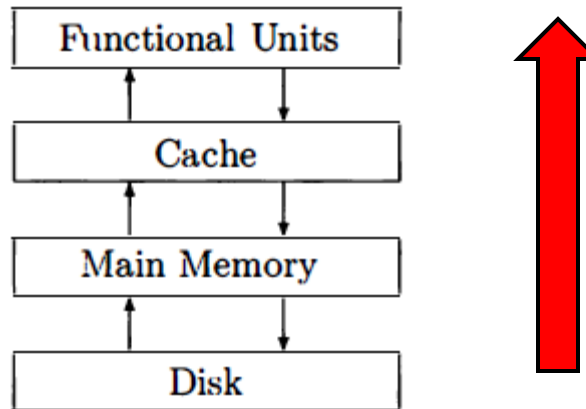
Outer Product Matrix Update

$$A = A + yx^T.$$

Load / Store Operations: (2+2n)

```
 $r_x \leftarrow x$   
 $r_y \leftarrow y$   
for  $j = 1:n$   
     $r_a \leftarrow A(:, j)$   
     $r_a = r_a + r_y r_x(j)$   
     $A(:, j) \leftarrow r_a$   
end
```

Memory Hierarchy



- ▶ Each level has a limited capacity.
- ▶ There is a cost associated with moving a data between two levels.
- ▶ Efficient Implementation should consider data flow between the various levels.

Parallel Matrix Multiplication

- ▶ Design of a parallel procedure begins with the breaking up of the given problem into smaller parts that exhibit a measure of INDEPENDENCE.
- ▶ Consider block matrices A, B, and C

$$C = \begin{bmatrix} C_{11} & \cdots & C_{1N} \\ \vdots & \ddots & \vdots \\ C_{M1} & \cdots & C_{MN} \end{bmatrix}, A = \begin{bmatrix} A_{11} & \cdots & A_{1R} \\ \vdots & \ddots & \vdots \\ A_{M1} & \cdots & A_{MR} \end{bmatrix}, B = \begin{bmatrix} B_{11} & \cdots & B_{1N} \\ \vdots & \ddots & \vdots \\ B_{R1} & \cdots & B_{RN} \end{bmatrix}$$

- ▶ The task is to compute:

$$\text{Task}(i, j): \quad C_{ij} = C_{ij} + \sum_{k=1}^R A_{ik} B_{kj}.$$



Data Motion in Parallel Multiplication

- ▶ In a parallel computing environment, the data that a processor needs can be “far away”, and if that is the case too often, then it is possible to lose the multiprocessor advantage.
- ▶ Time spent waiting for another processor to finish a task is TIME LOST.

