



@ODSC

ODSC EAST

Boston | May 3-5, 2017

Data Science with Spark:
Beyond the Basics

Adam Breindel

Instructor: Adam Breindel

LinkedIn: <https://www.linkedin.com/in/adbreind>

Email: adbreind@gmail.com



- ~20 years building systems for startups and large enterprises
- 10+ years teaching front- and back-end technology
- Fun big data projects...
 - Streaming neural net + decision tree fraud scoring
 - Realtime & offline analytics for banking
 - Music synchronization and licensing for networked jukeboxes
- Industries
 - Finance
 - Travel
 - Media / Entertainment

Course Objectives

- Use Spark to process dataset features
- Assemble processing, model-building, tuning and evaluation pipelines
- Understand strengths, patterns, mechanisms, and limitations of SparkML
- Perform operations that are not directly exposed in the Spark API
- Extend SparkML
 - custom feature processing
 - model algorithms

Not Objectives Today (due to time)

- Core Spark Job Execution, Cluster Configuration
- Data Science Basics, Fundamentals of Predictive Modeling
- "Parade of Algorithms"
 - We won't have lots of tiny examples showing every Spark ML API feature
 - Spark docs and examples folder contains those
- Math and CS of Distributed Algorithms (might cover a very small amount)

Today's Attendees - Spark Experience

This is an exciting workshop to run because there is minimal information ahead of time about attendees 😊

So let's take some informal polls

Approximate Schedule

Part 1 (Morning Session)

Welcome, Intro

Spark ML Patterns

Labs

Feature Engineering

NLP, Classification, Clustering

Part 2 (Afternoon)

Model-Parallel Cross-Validation w sklearn

Extending Spark:

Custom Feature Processing

Custom ML Algorithms

Gradient Descent

Deploying and Operationalizing Models

Files and Resources

Documents

- Slides, labs, data available at <http://tinyurl.com/odsc-spark>

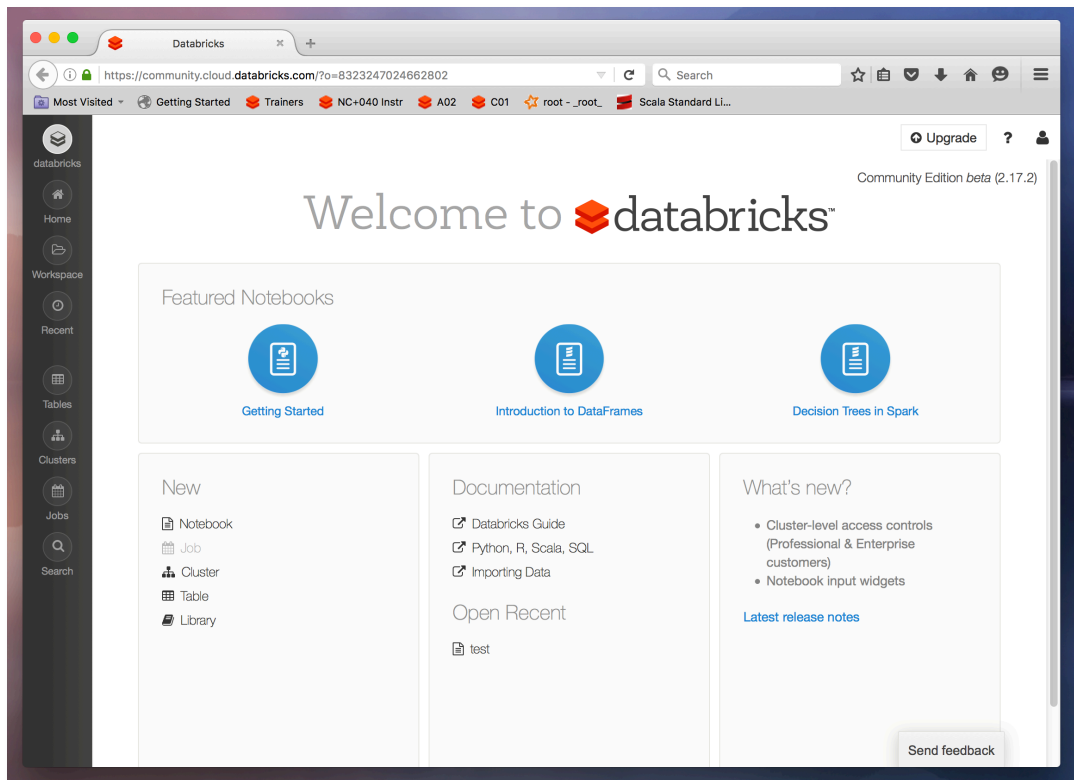
Databricks Account

- Create a Databricks CE account at <http://tinyurl.com/databricks-ce>
- Use a laptop with Firefox or Chrome (Internet Explorer / MS Edge not supported)

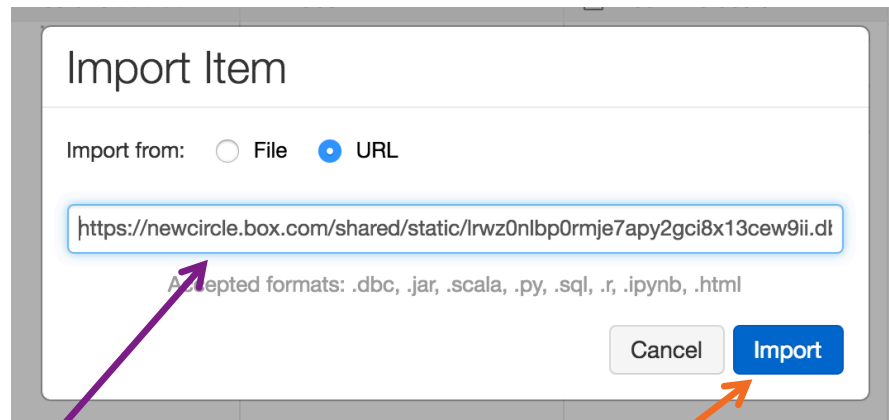
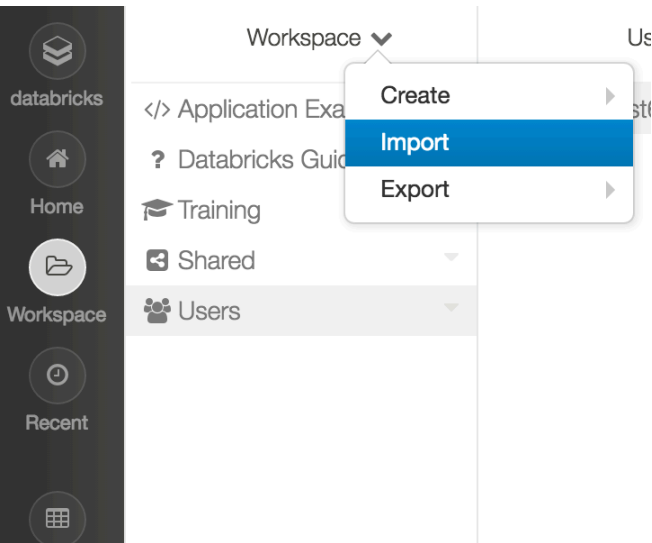
Getting Started

- Get ready type or copy this URL: <https://odsc-adbreind.c9users.io/labs.dbc>

Log in to Databricks




Import Labs into Databricks



Paste link from
previous page

Click import

Create a Spark Cluster!



Clusters

Active Clusters

+ Create Cluster

Name	Memory	Type	\$
------	--------	------	----

Terminated Clusters

Name	Memory	Type	\$
------	--------	------	----

Click to create cluster



Create Cluster

New Cluster

Cancel

Create Cluster

Cluster Name

LearnSpark

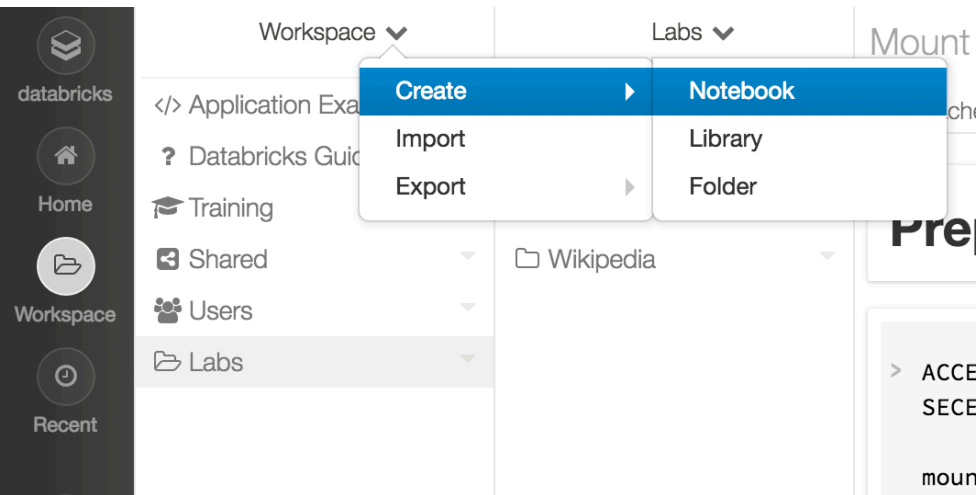
Apache Spark Version

Spark 1.6.1 (Hadoop 2)

Instance

- Enter a name
- Choose Spark-2.1.0-db3 (Scala 2.11)
- Click "Create Cluster" to finish

Create a New Notebook...



A 'Create Notebook' dialog box with the following fields:

- Name:
- Language:
- Cluster:

Buttons:

... and Test

- Key in `1 + 1` and press `SHIFT+ENTER` to see that the cluster is alive

Test (Python)

Attached: data ▼

View: Code ▼

File ▼

Run All

```
> 1 + 1
```

```
Out[1]: 2
```

```
Command took 0.07s
```

```
> %fs ls /mnt/training
```

▶ (2) Spark Jobs

path
dbfs:/mnt/training/bigrams/
dbfs:/mnt/training/data/
dbfs:/mnt/training/dataframes/

About Databricks CE

- Free, Community Edition for Learning/Exploring APIs
 - Not intended for Machine Learning at Scale
- Limited resources:
 - 6 GB RAM (container)
 - 0.88 cores (vCPU)
 - 8 threads (simulated cores)
 - will shutdown after 180 minutes idle time
 - start a new one ... notebooks/DBFS data does persist outside of cluster

About the Data

- Labs feature a variety of real datasets but...
 - Some are "small n" (fewer data points than we might like)
 - Some are "small p" (few predictors / dimensions)
 - Small n / small p can have different mathematical properties from large n / large p
 - Most are relatively clean in format and in terms of content
- If the real world were just like this ... it would be easier!
 - (Also we might not need Spark ☹)

What is *Spark*  ?

Spark provides...

- Distributed computation
- Data-parallel architecture
- Fault tolerance
- High scalability
- Data-locality-aware scheduling
- Load balancing



- Created by Matei Zaharia at UC Berkeley in 2009
 - Leverage cheap, fast local networking (10GB LAN)
 - Cheap(er) more abundant RAM
 - Less reliance on HDD (though disk still critical)
- Open Source (Apache License)
- Latest Release: v2.1.0 (December 2016)
- Repeatedly Rearchitected for Major (10x+) Perf Gains



Spark for ML: When and Why?

- Use Spark for ML iff you must train and/or predict on data sets too large for single-machine approaches
- Spark has a fairly narrow set of algorithms/models compared to, e.g., R
 - though it is extensible, and many 3rd-party libs are available
- Use Spark exactly when you have to scale out

mllib and spark.ml

mllib = original Spark ML API, based on RDDs

SparkML / MLPipelines / spark.ml = newer API, based on Dataset

- Spark 2.x: RDD-based APIs in spark.mllib have entered maintenance mode.
 - Primary ML API for Spark is now the DataFrame-based API in spark.ml
 - MLlib will still support the RDD-based API in spark.mllib with bug fixes
 - MLlib will not add new features to the RDD-based API
- Spark will add features to DF-based API to reach feature parity with RDD-based API
 - After reaching feature parity (~Spark 2.2), RDD-based API will be deprecated
- The RDD-based API is expected to be removed in Spark 3.0

Why switch to DataFrame-based API?

- DataFrames provide a more user-friendly API than RDDs
- Benefits of DataFrames include
 - Spark Datasources
 - SQL/DataFrame queries
 - Tungsten and Catalyst optimizations
 - Uniform APIs across languages.
- DataFrame-based API provides a uniform API across algorithms and multiple languages.
- DataFrames facilitate practical ML Pipelines, particularly feature transformations

Scala vs. Python

- This class is mostly in Scala... Why?
- We love Python too (and R)! But ... SparkML is built in Scala
 - API is designed to be almost identical in Python
 - but interacting with Scala gives more detailed feedback (e.g., types)
 - and direct access to every public part of the code
- If you extend Spark with new feature processors or models, using the standard pattern, you'll code them in Scala
 - and then maybe wrap them for Python access

Where are the Docs?

- Apache Spark ML Programming Guide site:
 - <http://spark.apache.org/docs/latest/ml-guide.html>
 - Tons of examples, multiple languages, explanations, links
- API Docs
 - <http://spark.apache.org/docs/latest/api/scala/index.html>
 - <http://spark.apache.org/docs/latest/api/python/pyspark.ml.html>
- New! "docs.databricks" site!
 - <http://docs.databricks.com/spark/latest/mllib/index.html>



SparkML API Patterns

Patterns Underlying ML

"Snap-together brick model"

- Encapsulation of processing
 - Transformer
 - Estimator
 - Pipeline
- Evaluation / Tuning
 - Evaluator
 - CrossValidator
 - ParamGridBuilder

Transformer

- Processes features
- Typically a "map" operation
 - E.g., Binarizer
- But can (occasionally) contain a reduce
 - E.g., OneHotEncoder
- Run by calling `aTransformer.transform(aDataframe)`
- Extend Spark by extending `UnaryTransformer` or `Transformer`

Estimator

- More complex feature processing, and/or model creation
- Typically one or many "reduce" operations
 - Builds (usually expensive and/or large) state
- Produces a Model (Transformer) to encapsulate state
 - Generate state & model by calling `anEstimator.fit(aDataFrame)`
 - Resulting model is a Transformer
- Extend Spark by creating a specific Model subclass and an Estimator that generates it

Pipeline

- Represents composition of
 - various Transformers' .transform methods
 - various Estimators' .fit and results' .transform methods
- "Point-free" operations
- Is itself an Estimator (supports composition)

Pipeline

Example Goal: rRun tf1 then tf2, then fit/apply est1, and fit est2

Instead of this ...

```
model = est2.fit(
    est1.fit(
        tf2.transform(tf1.transform(data))
    ).transform(
        tf2.transform(tf1.transform(data))
    )
)
```

We use this ...

```
model = Pipeline(stages=[tf1, tf2, est1, est2]).fit(data)
```

Uniform API

- Feature processors: Transformer or Estimator+Model
- ML Algorithms: Estimator
(> produce >)
- ML Models: Model
- Using any processor, algorithm, and performing tuning uses the same API!
 - Low cognitive load
 - "If you're bored, we're all winning!"
 - (because your brain is now free to work on the interesting hard stuff)



SparkML API Patterns (2)

Evaluator

- Calculates statistics on our models indicating
 - goodness-of-fit, explanation of variance
 - error quantities, precision/recall/etc.
- Generates 1 stat at a time
 - "mode-ful" switching of stat via setter
- Why? Designed for integration and for Spark, not just us
 - In particular, answers question "Which is better?" for tuning
- RegressionEvaluator, BinaryClassificationEvaluator, ...

ParamGridBuilder

- Helper to specify a grid of (hyper)params
 - Several params, chosen based on algorithm/model type
 - Several values for each param
 - Allows Spark to find/try every combination of values!

Parameter	Test Value 1	Test Value 2	Test Value 3	(etc.)
maxDepth	6	10	12	
maxBins	16	32	48	
(etc.)				

CrossValidator

- Performs K-fold cross-validation for...
 - combinations of param values from attached Param Grid
 - using attached Estimator (typically, algorithm or Pipeline)
 - and attached Evaluator (indicating better/worse results)

Example: 3-fold cross validation ... Divide training set in thirds, run 3 passes with the same data

